

anymcp

Any API's docs URL → a working MCP server. No spec required.

Every competitor needs a `swagger.json`.

We only need a **URL**.

Your AI agent hits a wall

Your agent needs to use **API X** — a payments provider, a data vendor, an internal service.

But there's **no MCP server**. And **no OpenAPI spec**. Just **HTML docs**.

Today that means: a developer reads the docs by hand, writes a wrapper, tests it, maintains it.

~1 week of engineering. Per API.

Every spec-based generator on the market stops right here — no spec, no server.

The long tail is enormous — and unreachable

The APIs with official MCP servers or clean OpenAPI specs are a **tiny minority**.

The other **thousands** publish only human-facing HTML docs:

- Regional **fintech** (India · UAE · SEA · LATAM)
- Niche **vertical SaaS**
- **Internal enterprise** APIs — never a public spec, ever

anymcp is the only tool that can turn *those* into agent-ready servers.


Live: a URL becomes a server

Watch.

1. Paste **CoinGecko's docs URL** — no spec, HTML only
2. `crawl` → `clean` → `extract` → `generate` (streams live in the browser)
3. A runnable **MCP server** — typed tools, auth, error handling
4. **Call it live** through the hosted gateway → real BTC/ETH prices

→ switch to the product at localhost:5173, then Claude Desktop

This is what the user sees

 **anymcp**

w2mcp

Build Call Install gateway localhost:8888

DOCS URL · OPENAPI · SWAGGER → MCP

Turn any API into a working MCP server

Have an OpenAPI/Swagger spec? We use it directly — complete and free. No spec? We read the HTML docs like a developer would. Either way your agent gets typed tools, auth, error handling, and a hosted endpoint.

docs URL or spec → crawl / parse → model → generate → MCP server



Spec-optional

No OpenAPI, no Swagger, no problem — we read the HTML docs the long tail only publishes.



Spec-native

Got OpenAPI/Swagger (JSON or YAML)? Converted directly — deterministic, complete, zero LLM.



Verified, not vibes

Every tool is live-probed against the real API before we ever call it "verified".



Hosted & encrypted

Multi-tenant gateway injects each tenant's credential per request; AES-256-GCM at rest.



Any MCP client

Drop it into Claude Desktop, Cursor, Claude Code, or run it over stdio. One config away.



Model-pluggable

The one AI step is comprehension only — swap Gemini, OpenAI, or Anthropic freely.

1 **Generate a server** from a docs URL or an OpenAPI/Swagger spec

☒ Render JS docs

Docs: Frankfurter

Docs: CoinGecko

Spec: httpbin

Generate →

2 **Call it live** hosted · multi-tenant · encrypted

Exactly what your agent does — a real call through the hosted gateway, credential injected server-side.

How it works

```
URL(s) → crawl → clean → extract → generate → MCP server
      (Playwright) (HTML→md) (LLM reads (typed tools, auth,
                               docs→model) shaping, errors)
```

5 of 6 stages are deterministic. The LLM only does *comprehension* — docs → a structured API model.

That makes the engine **model-pluggable** (Anthropic · Gemini · OpenAI) and provider-agnostic.

Have an OpenAPI/Swagger spec? We skip the LLM entirely and convert it directly — complete & free.

We're a **superset**: spec-optional, not spec-required.

No lock-in to one model vendor. Spec when you have it, docs when you don't.

Verified, not vibes

Generated tools can hallucinate endpoints. We refuse to call anything "verified" without a **real call**.

Three honest statuses on every server:

-  **live-verified** — endpoint actually called, returned 2xx + real data
-  **unverified-write** — write endpoint, flagged, checked against the doc example
-  **structurally-checked** — a guardrail, never claimed as proof

Trust is the product. This is the start of the moat.

Hosted, multi-tenant, secure

- Each API server runs as an **isolated subprocess** — a bad server can't crash the fleet
- Customers authenticate with an **anymcp key**; we inject **their** downstream credential **per request**
- Credentials **encrypted at rest** (AES-256-GCM) — never in the generated code, never on screen

Point us at your *internal* API docs behind the firewall → MCP servers for your whole stack.

This is the enterprise wedge: internal APIs never have specs, and switching costs are real.

The moat (what makes this un-clonable)

"URL-only" is the **wedge**. The **moat** is three compounding assets:

1. **Self-healing** — docs change, servers break; we monitor + auto-repair. *"Servers that stay working."*
2. **Verified** — a growing library of tested-green servers buyers can trust on day one.
3. **The registry** — the one directory with the long tail. More servers → more users → more feedback → better extraction. **A data + distribution flywheel.**

The generator gets the first 100 servers. The registry + self-healing keep everyone else out.

Market & model

Markets: US · UAE · Europe · India — the long-tail + internal-enterprise wedge is global.

Land → expand:

- *Land:* generate a server (self-serve, low friction)
- *Expand:* **hosting + monitoring + auto-healing + verification** — the recurring revenue *and* the moat, in one SKU

Don't fight Salesforce / ServiceNow / Xero (already have official MCP). Own everything they'll never build.

The ask

What we have: working engine · live-verified servers · hosted multi-tenant gateway · proven end-to-end.

Next 90 days: self-healing loop · public registry v1 · first 10 design-partner APIs.

Enterprise buyers → let's connect your internal APIs.

Investors → let's talk about owning the long tail of agent-ready APIs.

Kartik · 360-pm-team@awign.com